

UNIVERSITY OF BERGEN
DEPARTMENT OF INFORMATICS

Hidden parallelepiped in Hawk

Author: Eirik Djupvik Skjerve

Supervisors: Igor Aleksandrovich Semaev & Martin Feussner



UNIVERSITETET I BERGEN
Det matematisk-naturvitenskapelige fakultet

December, 2024

Abstract

Some abstract here

Acknowledgements

Thank you to some people

Eirik D. Skjerve

Tuesday 10th December, 2024

Contents

1	Introduction	1
1.1	Context and motivation	1
1.2	Objectives	2
1.3	Thesis outline	2
2	Background	3
2.1	Asymmetric cryptography	3
2.1.1	Digital signatures	3
2.1.2	Hash-Then-Sign	3
2.1.3	GGH	3
2.2	Linear algebra & lattices	3
2.2.1	Lattices	3
2.3	Probability theory	3
2.3.1	Distributions	3
3	Hawk and <i>Learning a parallelepiped</i>	4
3.1	Hawk	4
3.1.1	Simple Hawk	4
3.1.2	Key generation	4
3.1.3	Signature generation	4
3.1.4	Signature verification	4
3.2	Learning a parallelepiped	4
3.2.1	Solving the Hidden Parallelepiped Problem	5
3.2.2	HPP against normally distributed samples	6
4	Implementation	10
4.1	Implementation of Hawk	10
4.2	Implementation of HPP	10

5	Adapting HPP to Hawk	11
5.1	Covariance matrix secret key in Hawk	11
5.2	Secret Key Recovery	11
	Bibliography	12
A	Generated code	13

Chapter 1

Introduction

1.1 Context and motivation

Digital signatures are an integral part of secure communication today. The widely used DSA (Digital Signature Algorithm) and RSA (Rivest, Shamir, and Adleman) signature schemes are in peril due to the potential emergence of quantum computers which, theoretically, are able to break the hard problems DSA and RSA-sign are based upon. Whether practical quantum computers with these powers will emerge any time soon is debatable. However, measures against this potential looming threat has already begun. In 2016, NIST (National Institute of Standards and Technology) announced a process for selecting new standard schemes for Key Encapsulation Methods (abbr. KEMs) and digital signatures that are resilient against attacks based on quantum computers (source to NIST pages [here](#)). Many of the submissions to this process (including KRYSTALS-Dilithium which is to be standardized) are based on lattice problems that are believed to be hard to solve for both classical and quantum computers.

Cryptographic schemes based on lattice problems are not an entirely new phenomenon, however. The NTRU scheme and its signature counterpart NTRU-Sign, published in 2003 (source for this), is a digital signature scheme based on the hardness of the Closest Vector Problem. The original scheme was broken by Phong. Q. Nguyen & Oded Regev in 2006 [2], who showed that one can retrieve a secret key by observing enough signatures generated with one key. In other words, each signature leaks some information about

the secret key. A newer digital signature scheme submitted to NIST’s standardization process, Hawk [1], is somewhat similar to that of NTRU. One key difference is that Hawk signatures allegedly do not leak any information about the secret key in this manner due to the signature generation process.

1.2 Objectives

The objective for this thesis consists of two main parts:

- **Implementation of Hawk in Rust.** As the first part of this thesis I implement the Hawk digital signature scheme according to [1] in the Rust programming language. Implementing a scheme and its algorithms on ones own is a good way to learn how it works. I chose to implement it in Rust for the sake of learning this programming language as a bonus objective of the thesis. Moreover, having ones own version of an algorithm makes it easier to experiment, adjust and modify it to ones need. It would in any case be challenging to understand and work with dense, long, and complicated source code someone else has written. (For the Hawk teams source code and reference implementation see <https://github.com/hawk-sign>)
- **Cryptanalysis and experimentation.** The second part of this thesis is cryptanalysis of Hawk. The goal is to use the *Learning a parallelepiped* attack [2] and adjusting it to try and break Hawk. This requires both theoretical and practical work, and experiments will, like the Hawk implementation itself, be implemented in Rust.

1.3 Thesis outline

Chapter 2 will introduce important notions and mathematical background used in this thesis. Chapter 3 will introduce Hawk and the *Learning a Parallelepiped* attack, as well as discussing the implementation of these. In Chapter 4 our cryptanalysis of Hawk is presented. Chapter 5 will show results, and Chapter 6 will discuss future work.

Chapter 2

Background

2.1 Asymmetric cryptography

2.1.1 Digital signatures

2.1.2 Hash-Then-Sign

2.1.3 GGH

2.2 Linear algebra & lattices

2.2.1 Lattices

2.3 Probability theory

2.3.1 Distributions

Chapter 3

Hawk and *Learning a parallelepiped*

3.1 Hawk

In the following Hawk, the digital signature scheme, will be presented, as in [1]

3.1.1 Simple Hawk

Present simple sketch of keygen, sign and ver, as well as an example in dimension 2

3.1.2 Key generation

3.1.3 Signature generation

3.1.4 Signature verification

3.2 Learning a parallelepiped

The paper *Learning a Parallelepiped: Cryptanalysis of GGH and NTRU Signatures* by Phong Q. Nguyen and Oded Regev from 2006 introduced a method for breaking digital signature schemes based on the GGH scheme [2]. Essentially, by observing enough *message, signature* pairs generated by a secret key one can deduce this secret key. The attack broke the NTRU-Sign scheme by observing as little as 400 signatures.

3.2.1 Solving the Hidden Parallelepiped Problem

First we define an idealized version of both the problem to solve and the solution as proposed in [2]. Then we discuss the problem and solution in a practical context.

Hidden Parallelepiped Problem (HPP): Let $V = [\mathbf{v}_1, \dots, \mathbf{v}_n]$ be a secret $n \times n$ matrix and $V \in \mathcal{GL}_n(\mathbb{R})$. Define by $\mathcal{P}(V) = \{\sum_{i=1}^n x_i \mathbf{v}_i : x_i \in [-1, 1]\}$ a secret n -dimensional parallelepiped defined by V . Given a polynomial (in n) number of vector samples uniformly distributed over $\mathcal{P}(V)$, recover the rows \mathbf{v}_i of V .

We solve this problem in two main steps:

1. Transforming the hidden parallelepiped into a hidden hypercube
2. Learning the hypercube

Parallelepiped \rightarrow Hypercube The first step of the attack is to convert our hidden parallelepiped $\mathcal{P}(V)$ into a hidden hypercube $\mathcal{P}(C)$ with orthogonal basis vectors. Essentially, one moves each sampled point in accordance to a transformation matrix L computed the following way:

- Approximate $G \approx V^T V$ using our samples $\mathcal{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_u\}$
- Compute L such that $LL^T = G^{-1}$
- Then $C = VL$
- By multiplying our samples \mathcal{X} to the right by L , they are now uniformly distributed over the hidden hypercube $\mathcal{P}(C)$
- (—SHOW CALCULATIONS—)

Learning a Hypercube: The second step is to learn the hypercube. Given samples over $\mathcal{P}(C)$, we deduce the rows of the secret matrix C with the method described in Algorithm 1. After the rows are approximated, one can multiply the rows $\{\mathbf{c}_1, \dots, \mathbf{c}_2\}$ by L^{-1} such that we have $\{\mathbf{v}_1, \dots, \mathbf{v}_2\}$, and we are done.

Algorithm 1 Learning a Hypercube

Require: Descent parameter δ , samples \mathcal{X} uniformly distributed over $\mathcal{P}(C)$

Ensure: A row vector $\pm \mathbf{v}_i$ of C

Choose uniformly at random \mathbf{w} on the unit sphere of \mathbb{R}^n

loop

 Compute \mathbf{g} , an approximation of $\nabla \text{mom}_4(\mathbf{w})$

 Let $\mathbf{w}_{\text{new}} = \mathbf{w} - \delta \mathbf{g}$

 Place \mathbf{w}_{new} back on the unit sphere by dividing it by $\|\mathbf{w}_{\text{new}}\|$

if $\text{mom}_4(\mathbf{w}_{\text{new}}) \geq \text{mom}_4(\mathbf{w})$ **then** $\triangleright \text{mom}_4$ are approximated by samples

return \mathbf{w}

else

 Replace \mathbf{w} with \mathbf{w}_{new} and continue loop

end if

end loop

3.2.2 HPP against normally distributed samples

In the following, we see what happens to the computations the *Learning a parallelepiped* attack is based on if we replace the uniform distribution by a normal distribution. The key component and assumption of the *Learning a parallelepiped* attack is that the provided samples are distributed uniformly over $\mathcal{P}(V)$. Recall that $\mathcal{P}(V)$ is defined as $\{\sum_{i=1}^n x_i \mathbf{v}_i : x_i \in [-1, 1]\}$ where \mathbf{v}_i are rows of V and x_i is uniformly distributed over $[-1, 1]$ (generally one can take another interval than $[-1, 1]$ and do appropriate scaling). One runs into trouble if the sampled vectors are on the form $\mathbf{v} = \mathbf{x}V$ where \mathbf{x} follows a normal distribution, i.e. $x_i \sim \mathcal{N}(\mu, \sigma^2)$. Although one might be able to approximate the covariance matrix $V^t V$, one can not do a gradient descent based on the fourth moment given such samples; the fourth moment is constant since the samples form a hypersphere.

Adapting the definition of $\mathcal{P}(V)$ Firstly, the distribution $\mathcal{N}(\mu, \sigma^2)$ is defined over the interval $[-\infty, \infty]$, so it does not make sense to talk about samples "normally distributed over $\mathcal{P}(V)$ " without tweaking any definitions. Therefore, let $[-\eta, \eta]$ be a finite interval on which to consider a truncated normal distribution $\mathcal{N}_\eta(\mu, \sigma^2)$ such that $\int_{-\eta}^{\eta} f_X(x) dx = 1 - \delta$ for some negligibly small δ where $f_X(x)$ is the probability density function of $\mathcal{N}(\mu, \sigma^2)$. Now we consider $\mathcal{P}_\eta(V) = \{\sum_{i=1}^n x_i \mathbf{v}_i : x_i \in [-\eta, \eta]\}$ and proceed as in the original HPP with $\mathcal{P}_\eta(V)$ instead of $\mathcal{P}(V)$.

Approximating $V^t V$ Let $V \in \mathcal{GL}_n(\mathbb{R})$. Let \mathbf{v} be chosen from a truncated normal distribution $\mathcal{N}_\eta(0, \sigma^2)$ over $\mathcal{P}_\eta(V)$. Then $\lim_{\eta \rightarrow \infty} \mathbb{E}[\mathbf{v}^t \mathbf{v}] = V^t V \cdot \sigma^2$.

Proof. Let samples be on the form $\mathbf{v} = \mathbf{x}V$, where \mathbf{x} is a row vector where each element $x_i \sim \mathcal{N}_\eta(0, \sigma^2)$. Then $\mathbf{v}^t\mathbf{v} = (\mathbf{x}V)^t(\mathbf{x}V) = (V^t\mathbf{x}^t)(\mathbf{x}V) = V^t\mathbf{x}^t\mathbf{x}V$. Considering $\mathbb{E}[\mathbf{x}^t\mathbf{x}]$, we see that for $i \neq j$, $\mathbb{E}[x_i x_j] = \mathbb{E}[x_i]\mathbb{E}[x_j] = 0 \cdot 0 = 0$ due to independent random variables. For $i = j$, $\lim_{\eta \rightarrow \infty} \mathbb{E}[x_i^2] = \mathbb{V}[x_i] = \sigma^2$ since $\mathbb{V}[x_i] = \mathbb{E}[x_i^2] - \mathbb{E}[x_i]^2 = \mathbb{E}[x_i^2] - 0 = \sigma^2$. Therefore, $\lim_{\eta \rightarrow \infty} \mathbb{E}[\mathbf{x}^t\mathbf{x}] = I_n \cdot \sigma^2$, i.e. the matrix with σ^2 on the diagonal. Consequently, $\lim_{\eta \rightarrow \infty} \mathbf{v}^t\mathbf{v} = V^t \mathbb{E}[\mathbf{x}^t\mathbf{x}] V = V^t(I_n \cdot \sigma^2) V = (V^t V) \cdot \sigma^2$ and conversely $\lim_{\eta \rightarrow \infty} V^t V = (\mathbf{v}^t\mathbf{v})/\sigma^2$. \square

This means that we can in theory approximate the covariance matrix $V^t V$ by averaging over $\mathbf{v}^t\mathbf{v}$ and dividing by σ^2 . However, it is not immediately clear if one needs more samples for this approximation than in the original attack due to the difference in distributions. In addition, it is not clear exactly how accurate our approximation needs to be for the remaining parts of the attack to work.

Hypercube transformation Assume now that we know $V^t V$. Consider instead of $\mathcal{P}(V)$, $\mathcal{P}_\eta(V)$. Then by following part 1 of **Lemma 2** and its proof from [2] we can transform our hidden parallelepiped $\mathcal{P}_\eta(V)$ into $\mathcal{P}_\eta(C)$, a hidden hypercube, since this does not depend on what distribution the samples follow - it only assumes one knows $V^t V$. For completeness, by adapting the second part of **Lemma 2** to our case:

Proof. Let $\mathbf{v} = \mathbf{x}V$ where \mathbf{x} is normally distributed according to $\mathcal{N}_\eta(0, \sigma^2)$. Then samples \mathbf{v} are distributed according to $\mathcal{N}_\eta(0, \sigma^2)$ over $\mathcal{P}_\eta(V)$. It follows then $\mathbf{v}L = \mathbf{x}VL = \mathbf{x}C$ has a truncated uniform distribution over $\mathcal{P}_\eta(C)$. \square

Thus, we should also be able to map our normally distributed samples onto the hidden hypercube.

Learning a hypercube It is clear that samples uniformly over $\mathcal{P}_\eta(C)$ centered at the origin will form a hypersphere for which any orthogonal rotation leaves the sphere similar in shape. As a consequence, any measure of the fourth moment of one-dimensional projections is useless because every projection will yield the same result.

Analogous to [2] we compute the 2nd and 4th moment of $\mathcal{P}_\eta(V)$ over a vector $\mathbf{w} \in \mathbb{R}^n$. The k -th moment of $\mathcal{P}_\eta(V)$ over a vector \mathbf{w} is defined as $\text{mom}_{V,k} = \mathbb{E}[\langle \mathbf{u}, \mathbf{w} \rangle^k]$ where $\mathbf{u} = \sum_{i=1}^n x_i \mathbf{v}_i$ and $x_i \sim \mathcal{N}_\eta(0, \sigma^2)$. First we consider $\langle \mathbf{u}, \mathbf{w} \rangle = \langle \sum_{i=1}^n x_i \mathbf{v}_i, \mathbf{w} \rangle = \sum_{i=1}^n x_i \langle \mathbf{v}_i, \mathbf{w} \rangle$. Then for $k = 2$, $\mathbb{E}[(\sum_{i=1}^n x_i \langle \mathbf{v}_i, \mathbf{w} \rangle)^2] = \mathbb{E}[\sum_{i=1}^n \sum_{j=1}^n x_i x_j \langle \mathbf{v}_i, \mathbf{w} \rangle \langle \mathbf{v}_j, \mathbf{w} \rangle]$.

Due to independent random variables, $x_i x_j = 0$ when $i \neq j$, so we have $\mathbb{E}[\sum_{i=1}^n x_i^2 \langle \mathbf{v}_i, \mathbf{w} \rangle^2] = \sum_{i=1}^n \mathbb{E}[x_i^2] \langle \mathbf{v}_i, \mathbf{w} \rangle^2 = \sigma^2 \sum_{i=1}^n \langle \mathbf{v}_i, \mathbf{w} \rangle^2$ for sufficiently large n due to the well known result that $\mathbb{E}[x^2] = \sigma^2$ for $x \sim \mathcal{N}(0, \sigma^2)$. Thus, we end up with:

$$mom_{V,2}(\mathbf{w}) = \sigma^2 \mathbf{w}^T V^t V \mathbf{w} \quad (3.1)$$

We observe that if $V \in \mathcal{O}(\mathbb{R})$, $mom_{V,2}(\mathbf{w}) = \sigma^2 \|\mathbf{w}\|^2$

For $k = 4$:

$$\mathbb{E}[(\sum_{i=1}^n x_i \langle \mathbf{v}_i, \mathbf{w} \rangle)^4] = \mathbb{E}[\sum_{i=1}^n \sum_{j=1}^n \sum_{k=1}^n \sum_{l=1}^n x_i x_j x_k x_l \langle \mathbf{v}_i, \mathbf{w} \rangle \langle \mathbf{v}_j, \mathbf{w} \rangle \langle \mathbf{v}_k, \mathbf{w} \rangle \langle \mathbf{v}_l, \mathbf{w} \rangle]$$

We consider three cases for the indices i, j, k , and l :

1. **None equal:** if i, j, k , and l are different, the expression equals 0 due to independent random variables.
2. **All equal:** if $i = j = k = l$, then we have $\sum_{i=1}^n \mathbb{E}[x_i^4] \langle \mathbf{v}_i, \mathbf{w} \rangle^4$. A well known result for the normal distribution $\mathcal{N}(0, \sigma^2)$ is that $\mathbb{E}[x^4] = 3\sigma^4$.
3. **Pairwise equal:** if either
 - $i = j \neq k = l$
 - $i = l \neq j = k$
 - $i = k \neq j = l$

we have the following:

$$\sum_{i \neq j} \mathbb{E}[x_i^2 x_j^2] \langle \mathbf{v}_i, \mathbf{w} \rangle^2 \langle \mathbf{v}_j, \mathbf{w} \rangle^2$$

Since $\mathbb{E}[x_i^2 x_j^2] = \mathbb{E}[x_i^2] \mathbb{E}[x_j^2] = \sigma^4$ due to independent random variables, we have

$$\sigma^4 \sum_{i \neq j} \langle \mathbf{v}_i, \mathbf{w} \rangle^2 \langle \mathbf{v}_j, \mathbf{w} \rangle^2$$

Putting together the expressions above we have

$$mom_{V,4}(\mathbf{w}) = 3\sigma^4 \sum_{i=1}^n \langle \mathbf{v}_i, \mathbf{w} \rangle^4 + 3(\sigma^4 \sum_{i \neq j} \langle \mathbf{v}_i, \mathbf{w} \rangle^2 \langle \mathbf{v}_j, \mathbf{w} \rangle^2)$$

since there are three cases where indices pair up two and two. The final result becomes:

$$mom_{V,4}(\mathbf{w}) = 3\sigma^4 \left(\sum_{i=1}^n \langle \mathbf{v}_i, \mathbf{w} \rangle^4 + \sum_{i \neq j} \langle \mathbf{v}_i, \mathbf{w} \rangle^2 \langle \mathbf{v}_j, \mathbf{w} \rangle^2 \right) \quad (3.2)$$

Claim: If $V \in \mathcal{O}(\mathbb{R})$, and \mathbf{w} is on the unit sphere, $mom_{V,4}(\mathbf{w})$ is constant.

Proof. This can be shown by rewriting (3.2) as

$$\begin{aligned} mom_{V,4}(\mathbf{w}) &= 3\sigma^4 \left(\sum_{i=1}^n \langle \mathbf{v}_i, \mathbf{w} \rangle^4 + \sum_{i=1}^n \langle \mathbf{v}_i, \mathbf{w} \rangle^2 \sum_{j=1}^n \langle \mathbf{v}_j, \mathbf{w} \rangle^2 - \sum_{i=1}^n \langle \mathbf{v}_i, \mathbf{w} \rangle^4 \right) \\ mom_{V,4}(\mathbf{w}) &= 3\sigma^4 \left(\sum_{i=1}^n \langle \mathbf{v}_i, \mathbf{w} \rangle^2 \sum_{j=1}^n \langle \mathbf{v}_j, \mathbf{w} \rangle^2 \right) = 3\sigma^4 (\sigma^2 \|\mathbf{w}\|^2)^2 = 3\sigma^8 \end{aligned}$$

because $mom_{V,2}(\mathbf{w}) = \sum_{i=1}^n \langle \mathbf{v}_i, \mathbf{w} \rangle^2 = \sigma^2 \|\mathbf{w}\|^2$ when $V \in \mathcal{O}(\mathbb{R})$ and $\|\mathbf{w}\|^2 = 1$ when \mathbf{w} lies on the unit sphere. \square

In conclusion, if samples over the secret parallelepiped $\mathcal{P}_\eta(V)$ follow a continuous normal distribution, a gradient descent based on the fourth moment described in [2] is impossible because the fourth moment is constant over the unit sphere of \mathbb{R}^n .

The discrete Gaussian distribution We now do the same computations considering the discrete Gaussian distribution as described in [1]. It turns out that for the parameters used in Hawk, this distribution behaves like its corresponding Normal distribution, and consequently gradient descent is infeasible.

Consider the discrete Gaussian distribution as described in [1].

... Show definitions and stuff here ...

Then by computational results, as η grows large enough, $\mathbb{E}[x^2] = \sum_{-\eta}^{\eta} x^2 f_X(x)$ and $\mathbb{E}[x^4] = \sum_{-\eta}^{\eta} x^4 f_X(x)$ approaches σ^2 and $3\sigma^4$ respectively.

... Show table of computations here ...

Chapter 4

Implementation

Introduction to the implementation part of the thesis

4.1 Implementation of Hawk

Something something about the implementation of Hawk in Rust. Mentions of sampling, integer/float types, speed and comparison to Hawk team's C code?

4.2 Implementation of HPP

Something something about the implementation of HPP in Rust. Something about speedup/parallelization of gradient descent?

Chapter 5

Adapting HPP to Hawk

In this chapter we investigate the steps needed to possibly apply the Hidden Parallelepiped Problem to the Hawk digital signature scheme.

5.1 Covariance matrix secret key in Hawk

Nothing yet I'm afraid

5.2 Secret Key Recovery

Since \mathbf{x} follows some distribution close to some normal distribution, we hope that enough vectors \mathbf{w} will disclose some information about \mathbf{B}^{-1} . If we know \mathbf{B}^{-1} we know \mathbf{B} . This is the goal.

Bibliography

- [1] Joppe W. Bos, Olivier Bronchain, Léo Ducas, Serge Fehr, Yu-Hsuan Huang, Thomas Pornin, Eamonn W. Postlethwaite, Thomas Prest, Ludo N. Pulles, and Wessel van Woerden. Hawk. Technical report, NXP Semiconductors, Centrum Wiskunde & Informatica, Mathematical Institute at Leiden University, NCC Group, PQShield, Institut de Mathématiques de Bordeaux, September 2024.
URL: <https://hawk-sign.info/>.
- [2] Phong Q. Nguyen and Oded Regev. Learning a parallelepiped: Cryptanalysis of ggh and ntru signatures, 2009.

Appendix A

Generated code

Listing A.1: Source code of something

```
1 println!("Goodbye World");
```